Assignment 3 Presentation

Agenda

Data preprocessing and experimental set-up

Deep learning models

Training algorithm

Results

Data preprocessing and experimental set-up

Siamese Network

Generative Adversarial Network + CNN

Data Prep

1. Randomly Apply (50%)
    a. Random rotation (45 degrees)
    b. Random affine
        i. Translate / Shift
        ii. Scale / Zoom
        iii. Shear / Italicize
    c. Random flip (50%)
2. Normalize
3. Pairs as input

1. Normalize (mean=0.5, std=0.5)
2. GAN used to generate more training samples

Tools: pytorch*

Data selection

Images Background

Images Evaluation

Images Background Small 1

GAN generated training samples

Initial GAN

Good results but no way to tell which class is which.

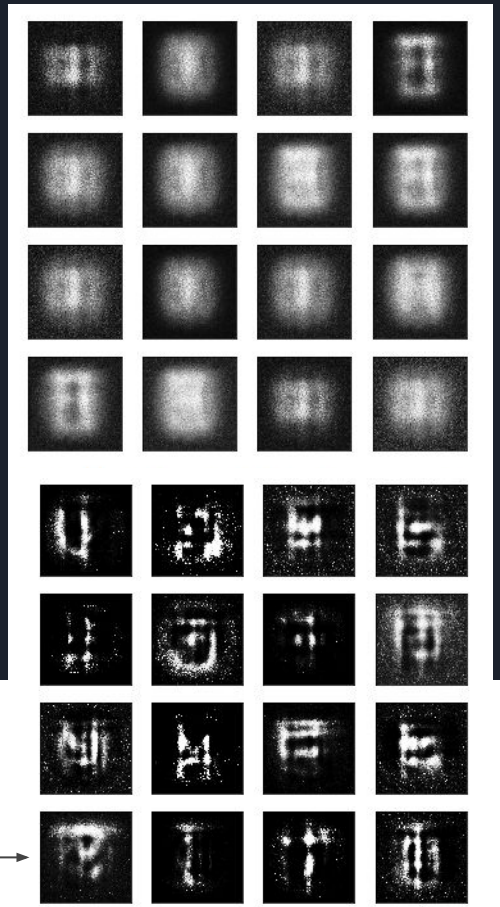

Look a dog

Using GAN to extend Data

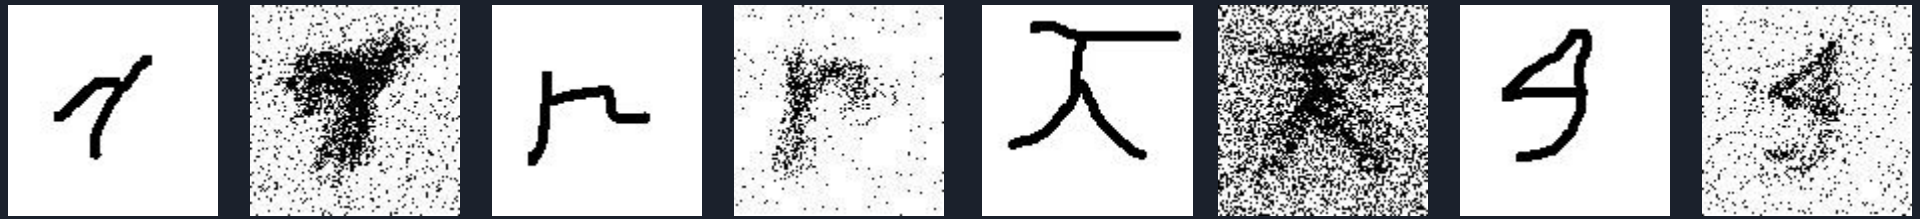

GAN is run for each character:

1. Noisy.. but is that good?
2. Takes a really long time

Model GAN

Discriminator:

Linear LeakyReLU Dropout

Linear LeakyReLU Dropout(0.3),

Linear LeakyReLU Dropout(0.3),

Linear Sigmoid

Tools: pytorch*

Generator:

Linear LeakyReLU

Linear LeakyReLU

Linear LeakyReLU

Linear Tanh

Tools: pytorch*

Model CNN

Feature / Conv

Conv2d BatchNorm2d ReLU

Conv2d BatchNorm2d ReLU MaxPool2d

Conv2d BatchNorm2d ReLU

Conv2d BatchNorm2d ReLU MaxPool2d

Classifier

Dropout Linear BatchNorm1d ReLU

Dropout Linear BatchNorm1d ReLU

Dropout Linear

Siamese Network Model

Conv1
Input: 1
Out: 6
Kernel: 5

ReLU
Max Pool
Kernel: 2
Stride: 2

Conv2
Input: 6
Out: 16
Kernel: 5

ReLU
Max Pool
Kernel: 2
Stride: 2

Conv1
Input: 1
Out: 6
Kernel: 5

ReLU
Max Pool
Kernel: 2
Stride: 2

Conv2
Input: 6
Out: 16
Kernel: 5

ReLU
Max Pool
Kernel: 2
Stride: 2

L1 Distance

8464

ReLU
Dropout: 50%

4232

1

Sigmoid

Training:

GAN was run over each character and after a certain number of epochs the output was saved to the dataset.

Once all the data was generated 2 CNN models of the same structure were run on the generated and original data.

Batch size of 100 was used. Learning rate of 0.0001.

Adam optimizer
Cross-Entropy as loss function

Results:

Extending the dataset with GAN output could be a valid way of reducing generalization error. As can be seen with an accuracy difference of ~7%.

Discuss performance:

Roadblock:

Thank you!